



Deep Learning Basics

Repo: <https://bit.ly/tf2-crash-course>

Why Learning?





Sheepdog or mop?

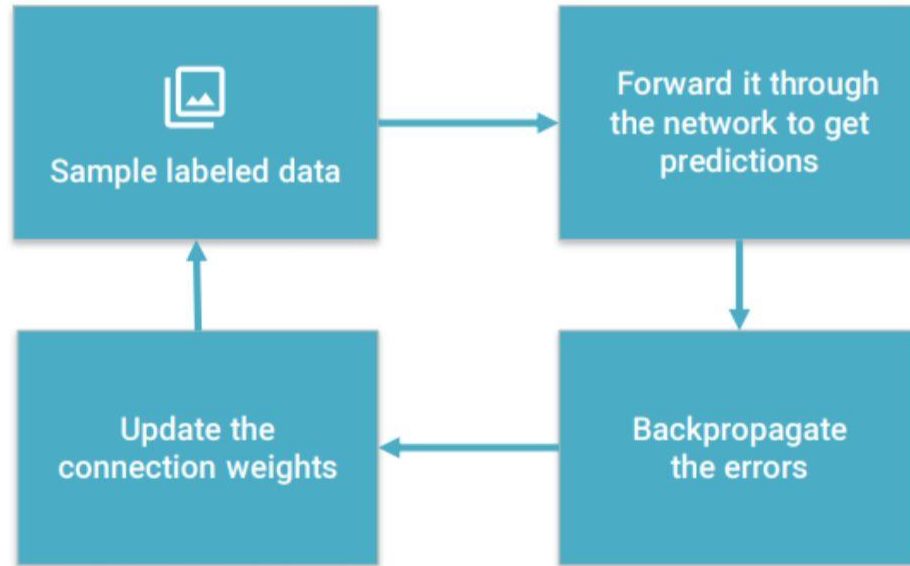


Chihuahua or muffin?

Neural Networks




Neural Networks - Training



Learns by generating an error signal that measures the difference between the predictions of the network and the desired values and then **using this error signal to change the weights** (or parameters) so that predictions get more accurate.

Training a Neural Network



- Perform parameter updates to **minimize** the loss (training objective)
- Typical flow involves:
 - Forward pass with the input going through various transformations
 - Compute the loss based on predictions and actuals
 - Compute gradients
 - Backpropagate gradients to update layer weights
- TensorFlow / PyTorch enables easy Automatic Differentiation

Custom Training Loops in TensorFlow

```
for epoch in range(epochs):
    print("\nStart of epoch %d" % (epoch,))

    # Iterate over the batches of the dataset.
    for step, (x_batch_train, y_batch_train) in enumerate(train_dataset):

        # Open a GradientTape to record the operations run
        # during the forward pass, which enables auto-differentiation.
        with tf.GradientTape() as tape:

            # Run the forward pass of the layer.
            # The operations that the layer applies
            # to its inputs are going to be recorded
            # on the GradientTape.
            logits = model(x_batch_train, training=True) # Logits for this minibatch

            # Compute the loss value for this minibatch.
            loss_value = loss_fn(y_batch_train, logits)


            # Use the gradient tape to automatically retrieve
            # the gradients of the trainable variables with respect to the loss.
            grads = tape.gradient(loss_value, model.trainable_weights)

            # Run one step of gradient descent by updating
            # the value of the variables to minimize the loss.
            optimizer.apply_gradients(zip(grads, model.trainable_weights))
```


- GradientTape records all relevant NN operations in the forward pass
- Hence easy to compute gradients in reverse order during the backward pass
- Useful to extract relevant gradients w.r.t the loss (used in adversarial attacks)

Deep Learning - ML on Steroids





Machine Learning



Deep Learning




Deep Learning Models - Multi-layered Neural Networks



Why is Deep Learning effective now?

New algorithms



More data



Software



theano



PYTORCH



Faster compute engines



Deep Learning Model Architectures

1 Convolutional Neural Networks (CNNs)

Used extensively in computer vision problems with image, video. Can also be used for audio and text

2 Recurrent Neural Networks

Good for sequential data, used for time series forecasting and NLP problems

3 Long Short Term Memory Networks (LSTMs)

Can remember longer sequences of data and better than RNNs

4 Gated Recurrent Units (GRUs)

Can remember longer sequences of data and faster than LSTMs

5 Bi-directional Models

Processes sequences of data in both directions for capturing better contextual information

6 Auto-encoders

Learns efficient latent data representations in lower dimensions using unsupervised learning

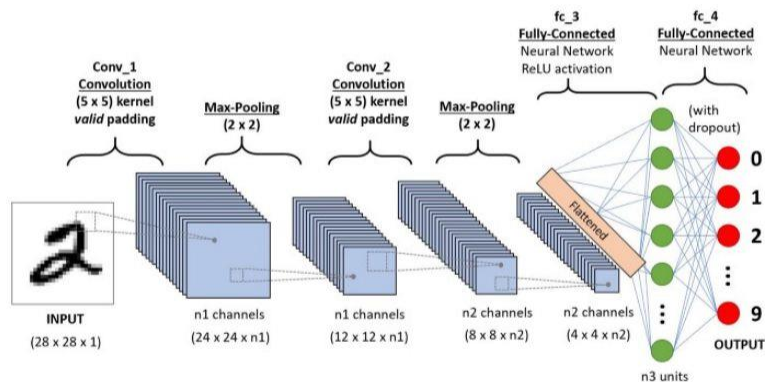
7 Encoder-Decoder Models

Takes in a sequence of data and generates a sequence of data as output

8 Transformer Models

Stack of encoder-decoder models used for language modeling and can be tuned for different NLP tasks

Convolutional Neural Networks (CNNs)



- CNNs have a layered architecture of several layers to learn hierarchical spatio-temporal features
- Convolution Layers use convolution filters to build feature maps (feature extraction)
- Pooling Layers help in reducing dimensionality after convolutions (compression)
- Non-linear activation functions are applied in the network as usual
- Dropout or BatchNormalization Layers may be used to prevent model overfitting
- FC Layers in final stages help with flattening and prediction

Convolutional Neural Networks (CNNs)

